

A Cloud Based Lab Environment for Network Security Research

Franklin E. Diaz

December 27, 2022

Abstract

The goal of this document is to detail the major components of the “Cloudlab” network security lab in [Google Cloud](#). The reasoning behind including certain components and design elements is explained. The relationship between security and the lab are detailed to highlight purpose. Security learning objectives are realized at every step of the process in the creation and use of this lab. The result is a widely scoped and highly flexible security learning environment for data scientists and engineers with focus on development, test, operations, and continuous pipelines.

[Click here to download latest version.](#)

Cloud-native technologies empower organizations to build and run scalable applications in modern, dynamic environments such as public, private, and hybrid clouds.

[Cloud Native Computing Foundation](#)

1 A Cloud Native Security Lab

There is a proliferation of new and constantly evolving Public Cloud providers, tools, tool chains, and whole “cloud based” development and operations ecosystems. There is a need to understand and adapt to the “GitOps” paradigm and the business opportunities this fundamental shift presents. There is no disputing the fact that this shift is well underway[4].

The “Cloudlab” is a private lab environment for security research, testing, and training. The lab is “Cloud Native” in the sense that it adheres to [GitOps practices](#). The GitOps methodology facilitates storage, review, and maintenance of the lab Infrastructure as Code. The lab is currently comprised of two main Pipelines. The first of these is dedicated to Continuous Integration. The second pipeline facilitates [Continuous Machine Learning](#), or CML[1].

As security practitioners, it is important for us to understand the infrastructure and applications being built in the public clouds as a first step to making things more secure.

2 Project Goals

There are several goals related to this project. These goals overlap and interconnect at times. The overarching goal is to understand the composition, management, and weaknesses of the items listed here.

1. Provisioning [Palo Alto Networks CN-Series firewall](#) products, integrating them with [Calico](#) and protecting containerized workloads (Kubernetes “pods”).
2. Research containerized deployments, workloads and security.
3. Demonstrate “[automation bots](#)” and [how they interact with a GitHub repository](#). Extending a revision control platform (GitHub in this case) is rather common in large organizations.
4. Integration of [Bridgecrew “Security as Code” tooling](#) with GitHub repositories.
5. Develop and demonstrate “serverless cloud function” expertise (GCP Cloud Functions in this case).

6. Develop and use a cloud native Continuous Integration build pipeline. The output of this pipeline is a Docker image that is stored in gcr.io. These images include a fully contained set of tools, documentation and Terraform code for customer deployments.
7. Demonstrate Policy as Code concepts using [Terratest](#) and [Kyverno](#).

3 Learning Objectives

Building and operating this lab is meant to foster learning. It may be beneficial for engineers to focus on some or all of the following list of topics.

1. API Endpoints
2. Deployment and operation of CN Series firewalls.
3. Kubernetes role based access control (RBAC) and security.
4. Developing serverless functions in Python.
5. ML/AI pipeline, containerized workloads, and their security.
6. Containers and container security.
7. Automation.
8. Testing perspectives including Policy as Code, Security as Code, etc.
9. CI/CD Pipelines and security.

Training materials derived from lab construction and operation are easily within reach of PS engineers.

3.1 API Endpoints

Construction of, and interaction with API endpoints is a good space for exploring potential security issues in cloud native environments. The API is foundational in how applications interact between all sorts of environments including serverless and Kubernetes.

3.2 Containers and Container Security

3.3 Automation

4 Lab Configuration Details

The configuration files and code used to build the lab are stored on GitHub. The lab is running on Google Cloud, one of the “big three” public cloud providers.

A Cloud Native Security Lab

The lab is intended to serve as a proof of concept as well as provide a teaching and training platform.

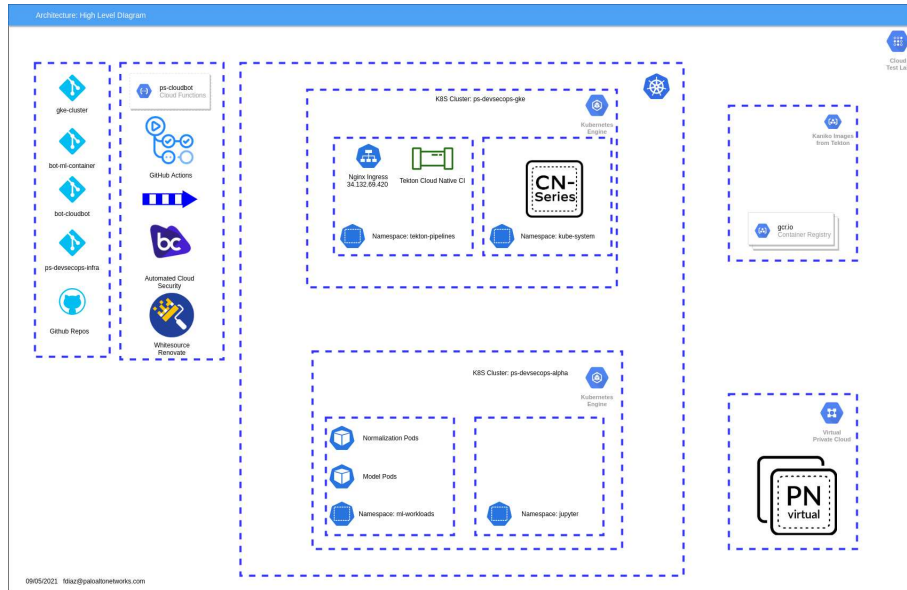


Figure 1: High Level Lab Design Diagram

A detailed description of the five main blocks seen in figure 1 follows.

1 [GitHub Repositories]

The code base for the lab is broken down into several GitHub repositories, more or less around the functional area.

Repo Name	Purpose
bot-cloudbot	A custom Python 3.9 GCP Cloud Function for GitHub pull request task automation.
bot-ml-container	An experimental containerized machine learning model.
gke-cluster	Infra as Code files for GKE cluster, YAML files for Tekton CI pipeline . CN Series firewall nodes.
ps-containerizer	An “invisible shim” with a Docker image for each “public cloud” VM-Series Terraform module development repo. Allows PRs to be ingested into Tekton CI pipeline without any integrations with the source repository.
ps-devsecops-alpha	IaC files for Alpha K8s cluster.
python-project-template	Python project template for writing serverless code in AWS Lambda and GCP cloud functions.

Table 1: Project Codebase - GitHub Repositories

2 [GitHub Actions]

GitHub repositories that have been “on-boarded” to the project have certain “actions” included.

1. [Bridgecrew](#) is used to scan all commits to all open pull requests.
2. Whitesource Renovate is used to track keep project dependencies up to date and secure.
3. Another GitHub action defines the parameters of the GCP project and helps the “ps-cloudbot” with pull request maintenance tasks.

Note that there is also a [webhook](#) to make the CI pipeline aware of each commit and kick off a test and build cycle.

3 **[Kubernetes Clusters]** Two clusters are deployed with the Google Kubernetes Engine (GKE). The “gke” cluster hosts the [Tekton CI pipeline](#). The “alpha” cluster is used for machine learning experimentation.

Pipeline runs can be viewed and managed through a graphical interface that is well suited for development teams. There is also the ability to manage pipeline runs and their requisite tasks using standard command line tooling.

4.1: Example pipeline command

```
franklin ~/gke: tkn pr ls
```

NAME	STARTED	DURATION	STATUS
gh-pr-run-ncwrr	1 hour ago	1 minute	Succeeded
gh-pr-run-vhrvs	2 hours ago	1 minute	Succeeded
gh-pr-run-q6szq	3 hours ago	1 minute	Succeeded
gh-pr-run-8vn22	6 hours ago	1 minute	Succeeded
gh-pr-run-h7twc	14 hours ago	17 seconds	Failed
gh-pr-run-74wm7	21 hours ago	1 minute	Succeeded
gh-pr-run-tqlfg	1 day ago	1 minute	Succeeded
gh-pr-run-xs52f	1 day ago	1 minute	Succeeded
gh-pr-run-xtdz6	1 day ago	1 minute	Succeeded
gh-pr-run-w2zn7	1 day ago	1 minute	Succeeded

4 [GCR Private Container Repository]

To date there are about 15 GitHub repositories that are integrated with the CI pipeline outlined in this paper. Each commit to a pull request causes the CI pipeline to generate a Docker image. These Docker images are [stored in a private GCR location](#).

5 [Panorama Management]

Panorama is a key component in deployment and maintenance of Kubernetes clusters and CN Series firewalls. Currently we have two virtual Panorama devices deployed in a high availability (HA) configuration.

Palo Alto Networks has historically maintained serverless functions, for example in AWS Lambda, for firewall “auto-scaling” tasks. This has since been replaced by a set of official Panorama “plug-ins” that can be downloaded to PanOS devices. Lambda and Cloud functions are still frequently used to augment the capabilities of this new family of plug-ins.

5 Serverless Functions

Google Cloud offers [Cloud Functions](#), among other serverless computing offerings. This is a good set of patterns to learn and comes up often in security infrastructure work since organizations can use it to run jobs at a lower cost. Securing the webhooks and connection between cloud functions and the VPC the GKE cluster is in is an important factor in deployments. Making XML or other sorts of API calls is often used to pass data between management platforms, ticketing systems, custom applications, and so on.

[Here is a full working example](#) of how to integrate Github, Cloud Functions, and GKE clusters in the manner described in this section.

5.1 Github Webhook

The Github webhook can be combined with a set of GCP credentials stored as a “repository secret” in a repository. Combined with [a Github action as a triggering mechanism](#), we have a powerful and flexible method to combine systems into a more intricate Continuous Integration and testing pipeline.

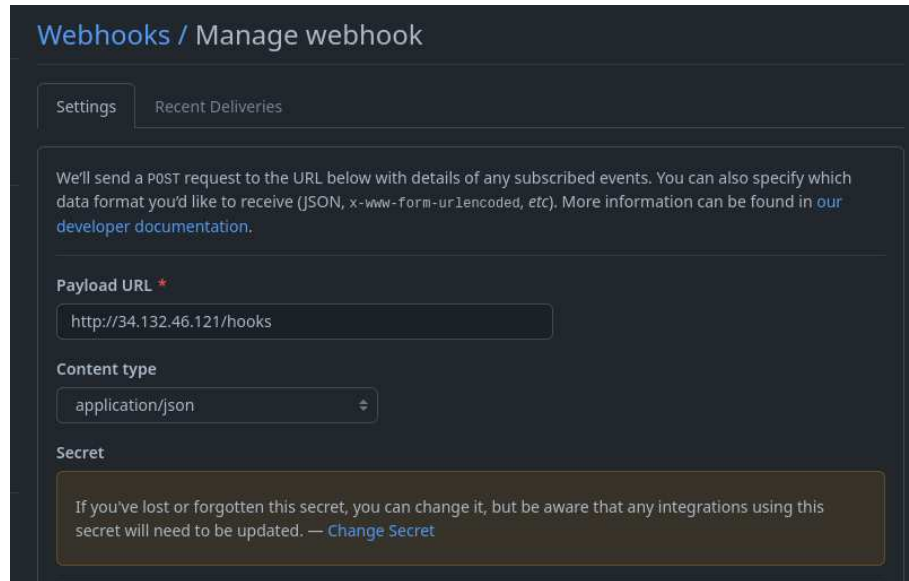


Figure 2: Webhook configuration settings in Github




Figure 3: Check only this box in the webhook settings

5.2 Connecting Serverless to GKE

A network peering is created between the Cloud Function and the VPC that the GKE cluster resides in. This allows the serverless function application to make calls to a containerized deployment, such as a Node application.

5.1: Example YAML for a cloudbot service

```
apiVersion: v1
kind: Service
metadata:
  name: cloudbot-service
  namespace: ci-build
  annotations:
    cloud.google.com/load-balancer-type: "Internal"
  labels:
    app: cloudbot-service
spec:
  type: LoadBalancer
  selector:
    app: cloudbot
  ports:
    - port: 80
      targetPort: 8089
      protocol: TCP
```


5.2: Example deployment YAML for a cloudbot

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: cloudbot
  name: cloudbot-deployment
  namespace: ci-build
spec:
  replicas: 3
  selector:
    matchLabels:
      app: cloudbot
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: cloudbot
    spec:
      nodeSelector:
        env: build
      containers:
        - name: cloudbot
          image: gcr.io/gcp-gcs-pso/build-pod
          imagePullPolicy: Always
          volumeMounts:
            # name should match from volumes section
            - name: nfs-volume-1
              mountPath: "/data"
      volumes:
        - name: nfs-volume-1
          persistentVolumeClaim:
            claimName: nfs-pvc
```

5.3: Example Dockerfile for a cloudbot pod in GKE

```
# THIS IS AN EXAMPLE Dockerfile

FROM node:22-alpine

LABEL maintainer="Franklin D <devsecfranklin@duck.com>"
LABEL org.opencontainers.image.source=https://github.com/devsecfranklin/paper-cloud-
LABEL org.opencontainers.image.description="Cloudlab Paper"
LABEL org.opencontainers.image.licenses=MIT

ENV NODE_ENV=production

COPY ["app.js", "package.json", "package-lock.json", "./"]

RUN npm install --production

COPY . .

ENTRYPOINT ["node", "app.js"]
```

6 Continuous Integration

Keeping internal and external build pipelines secure, as well as scanning work products that traverse these pipelines is highly desirable. To that end, a fully operational Cloud Native Continuous Integration pipeline has been implemented in the lab.

The pipeline can ingest a code base from a public repository, or a private repository with proper credentials. The “containerizer” repo demonstrates the ability to collect files from a repository, bundle it with requisite command line tools, test cases, and other necessities, and ship the image to the gcr.io container registry at the conclusion of a successful pipeline run. A pipeline run refers to a set of test cases managed and executed by [Tekton](#).

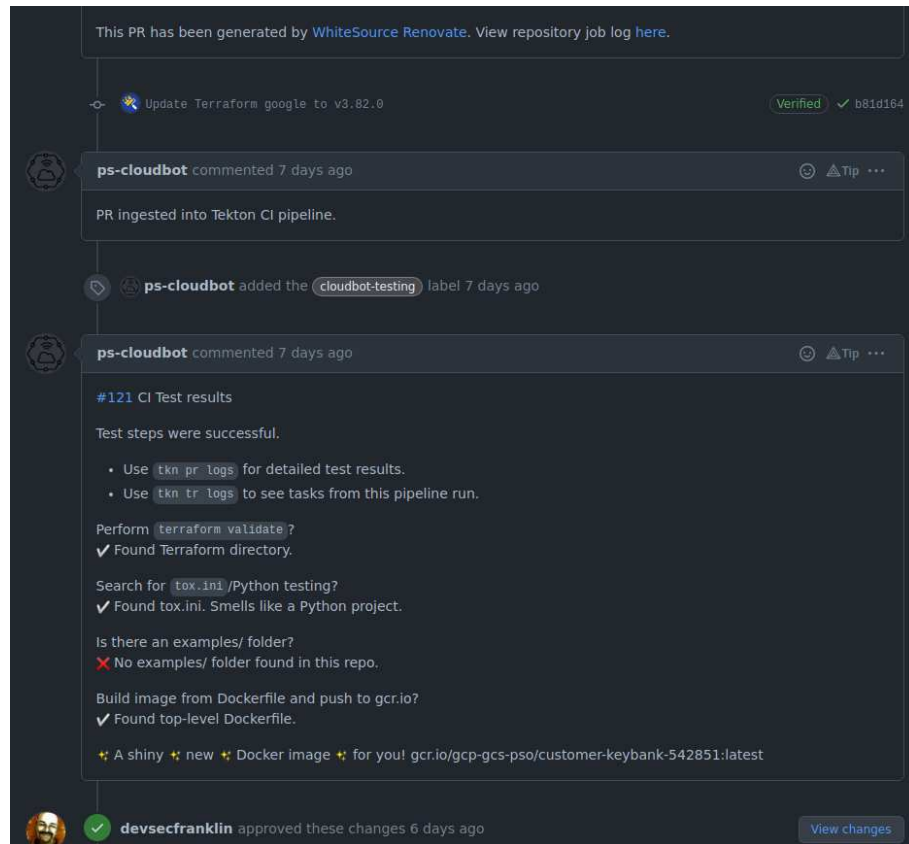


Figure 4: Tekton automated pipeline run results in a GitHub comment

Consider figure 4. The “Renovate” bot detects an out of date or insecure dependency. A pull request is opened by the Renovate bot. Next, “ps-cloudbot” GCP cloud function is notified of the new pull request via webhook. The second bot places a label on the pull request to indicate it is performing administrative actions on the pull request. The cloud function might perform other actions such as assigning the pull request to a certain user, adding certain users as reviewers, providing documentation, and so on. In this example, the cloud function adds a comment to the pull request to inform project members that it has been accepted by the CI pipeline. A set of test cases based on certain technology or functional area is executed and the results are returned to the pull request in a second comment. There is also a link to the container image in the container registry. Although the pull request is merged manually in this instance, the workflow could be modified to approve and merge the pull request with no human interaction whatsoever.



Figure 5: Bridgecrew integration with pull request automation

In addition to the ability to scan for dependencies which are out of date or have vulnerabilities of note, we can perform automated security scanning of the code base as seen in figure 5. As before, the Renovate bot has detected the use of an out of date Docker base image for Golang. Because the Dockerfile is updated as part of this pull request, Bridgecrew scans the file and triggers on misconfigurations that may lead to security issues. The results of the scans and the issues found are noted in the pull request comments. Automated remediation and merging of these issues may be possible in some cases.

7 Summary

Future goals for the lab project might include, but are not limited to the following list.

1. Orchestration of Multi-Cloud infrastructure builds and deployments, using [Crossplane](#) for example.
2. Understanding infrastructure as code in languages besides HCL (aka Terraform). One example of this is [Pulumi](#).
3. “Red teaming” our own lab, conducting dynamic application security testing (DAST) exercises.

Revision History

Revision	Date	Author(s)	Description
v0.1	Sept. 7th, 2021	Franklin Diaz	Initial Draft
v0.2	January 29th, 2022	Franklin Diaz	incorporate feedback
v0.3	January 29th, 2022	Franklin Diaz	pin release for ORCID
v0.4	December 27th, 2022	Franklin Diaz	GKE updates

References

- [1] Ethem Alpaydin. *Machine Learning: The New AI*. MIT Press, 2016.
- [2] Brendan Burns. *Kubernetes best practices : blueprints for building successful applications on Kubernetes*. O'Reilly Media, Sebastopol, CA, 2019.
- [3] Gene Kim. *The DevOps handbook : how to create world-class agility, reliability, & security in technology organizations*. IT Revolution Press, LLC, Portland, OR, 2016.
- [4] Christy Pettey. Cloud shift impacts all it markets. <https://www.gartner.com/smarterwithgartner/cloud-shift-impacts-all-it-markets/>, 2020. [Online; accessed 2021-09-05].
- [5] Dan Sullivan. *Official Google Professional Cloud Architect : study guide*. Sybex, a Wiley brand, Hoboken, NJ, 2020.